

OttoBot x Filipinohomes: MVP PRD

Version: 2.0 (2026-10-06). Supersedes `.planning-archive/REQUIREMENTS.md` (v1, generic SMB spec, kept read-only). **Owners:** Jiv Tuban (agent, channels, compliance) and Gabe (surfaces, auth, data). Split in section 7. **Evidence:** every external fact below cites `docs/research/2026-10-06-filipinhomes-deep-research.md` (C1 to C9). Anything not in its CONFIRMED list is marked **[unverified]** or listed as an open question. No price, size or response-time number is used unless confirmed.

1. Brief (one screen)

What: Pilot OttoBot as a Taglish first-responder and lead qualifier for Filipinohomes. A buyer inquiry gets a reply in seconds, is qualified on the fields a real-estate agent actually needs, and is handed to the right human agent with a summary and a proposed site-viewing slot.

Why Filipinohomes: they are a beta tester and potential partner, and their own agent directory publishes an "Inquiry Reply" time per agent. The fastest agents sit at about 10 to 48 minutes and the list falls to hours below that (C4, live page checked 2026-10-06, 12 of about 1,737 agents, method unstated). That is a public, measurable gap and our beta metric.

What we ship: one tenant (Filipinhomes), text only, three lead sources, two reply channels, handoff to human agents, an ops dashboard that proves the reply-time win.

What we do not ship: voice calls, a second customer, billing, agency multi-account (old Phase 7), the autoresearch loop (old Phase 8), calendar integrations, a built-in CRM.

Done means: a real lead from a Filipinohomes source gets an AI first reply under 60 seconds (target, not researched), is qualified, and lands with an assigned agent who confirms a viewing, with every step visible in the dashboard and a consent and opt-out trail on record.

2. Who Filipinohomes is (verified vs not)

Fact	Status
filipinhomes.com is a Cebu-based real-estate company led by founder and CEO Anthony Gerard O. Leuterio, active July 2026	Confirmed (C1). Source is their own news site. Confirm it is your contact's company.
Hybrid: nationwide network of property specialists plus listings search and an agent directory. No legal entity named on the about page	Confirmed medium (C2). Portal vs brokerage is contested.
Size: "144 branches, 35,000 agents" (their news) vs "52 offices in 40 cities" (press, per verifier)	Unverified, conflicting (C3). Do not quote either.

Fact	Status
Per-agent "Inquiry Reply" time and a "Sort by Fastest reply"	Confirmed medium (C4).
Lead routing, tech stack, integrations, volume, revenue model, existing AI tooling	Unknown. Needs the beta contact (section 9, Q1 to Q8).

Do not assume the agent base is a long tail of solo agents; that claim was refuted 0-3.

3. Users and jobs

User	Job	MVP surface
Buyer or inquirer (Taglish, sometimes English, Bisaya [unverified])	Get an answer now, without waiting for an agent	Messenger and SMS
Filipinohomes agent	Receive a qualified lead with context, confirm a viewing	Mobile app + push, web inbox
Filipinohomes ops/admin	See reply time, funnel, who is slow; control assignment	Web dashboard
Jiv and Gabe	Prove the win with numbers and keep the pilot safe	Metrics, audit log

4. Current state (verified in the repo 2026-10-06)

- 144 backend tests pass on `feat/neon-clerk-migration`. Phases 1 to 6 of v1 shipped.
- Real-estate persona exists but is thin: `agent/prompts/real_estate.j2` is 14 lines, with no qualification fields.
- Channels exist: Semaphore SMS (`/webhook/sms`, `api/main.py:556`), Messenger (`/webhook/facebook`, `:583`, `:600`), lead upload (`:633`), lead form (`:695`).
- **Single-tenant by environment variable:** all three webhooks read `BUSINESS_ID` and `DEFAULT_INDUSTRY` from env (`api/main.py:567`, `:609`, `:705`). One deployment serves one business.
- **Lead form handler will not work on real Meta traffic.** `/webhook/lead-form` reads `value.leads[].field_data` from the webhook body (`api/main.py:715`). Meta's leadgen webhook delivers a `leadgen_id`, and the lead data is fetched with a separate retrieval call (Meta doc, updated 2026-05-21, C8). Needs rework.
- **Escalation targets one owner** (push to a single token plus Resend email). Filipinohomes needs routing to a specific agent.
- AI disclosure: `agent/prompts/base.j2:2-8` discloses AI on inbound conversations and skips it on the outbound first message. This is the behavior fixed on 2026-07-27, and it is a decision to revisit (section 8, D3).
- Migration leftovers still on Supabase: `mobile/lib/supabase.ts` (stub), `mobile/app/(tabs)/settings.tsx`, `frontend/src/OnboardingWizard.tsx`, `tests/conftest.py`. Tasks

7 to 10 of docs/superpowers/plans/2026-07-29-supabase-to-neon-clerk.md cover them.

- Login design: the restrained redesign is not confirmed wired into mobile/app/login.tsx . Re-verify before Gabe starts G6.
- Design system (signal green, accent discipline) is defined in docs/design/guardrails.md and applies to every new screen.

5. MVP scope

P0 (pilot cannot start without these)

ID	Requirement	Acceptance
RE-01	Real-estate qualification: capture budget, preferred location, property type, pre-selling vs ready-for-occupancy, financing route (cash, bank, Pag-IBIG, in-house), OFW status, timeline, viewing availability	Fields stored per lead, shown in inbox, agent asks only what is missing, no field asked twice
RE-02	Taglish and English replies that match the lead's register; "po/opo" mirroring kept	Eval set of 30 scripted conversations, pass rate recorded per model
LEAD-01	Lead sources: Facebook Lead Ads (leadgen_id plus Graph fetch), inbound Messenger, inbound SMS, CSV upload	Each creates one deduped lead with source set
CH-01	Tenant routing replaces BUSINESS_ID env: webhook resolves business from page id, number, or token	Two test tenants cannot see each other's leads
CH-02	Messenger and SMS send/receive with Meta messaging-window rules respected	Out-of-window sends are blocked and logged, not silently dropped
HAND-01	Handoff to the assigned human agent with summary, qualification card and proposed viewing slots	Agent gets push and sees the card; lead is told who follows up
HAND-02	Agent roster and assignment rule (round robin or by location)	Admin imports agents from CSV and edits the rule
CONS-01	Consent and notice: privacy notice in the first message or form, STOP/opt-out honored on every channel, audit log of consent and opt-out	Opt-out stops all sends within one turn; log is queryable
MET-01	Metrics events: inquiry received, first AI reply, qualified, handed off, human first reply, viewing confirmed	Dashboard shows median first reply (AI vs human), funnel, per-agent table
AUTH-01	Neon plus Clerk cutover finished, Supabase fully removed	git grep -i supabase returns docs and archive only
UI-01	Web inbox, conversation view, agent roster, pilot dashboard; mobile pipeline and push	Follow docs/design/guardrails.md ; reduced-motion safe

P1 (build if P0 is green with time left)

- Viber Business sender (cost model in C5; needs a Philippines-registered sending entity or the rate is about 14x to 29x higher).
- Bisaya support, tested on a separate eval set.
- English admin copy pass for CRM (Task 10), keep agent-facing Tagalog.
- Weekly pilot report export (CSV).

Out of scope for the MVP

Voice calling (unresearched, text first), WhatsApp, second tenant, billing, calendar sync, listing-aware recommendations from the portal catalog, auto-optimization loop, email channel.

6. Success metrics (pilot)

Baseline first, targets second. Targets below are design goals, not researched benchmarks.

Metric	Baseline source	Pilot target
Median first reply time	Filipinohomes data, to be requested (public per-agent figures are 10 minutes to 15 hours, method unknown)	AI first reply under 60 s, measured by MET-01
Inquiry to qualified rate	Not known	Report only in week 1, set target after
Qualified to viewing confirmed	Not known	Report only in week 1, set target after
Opt-out and complaint rate	None	Zero un-honored opt-outs
Agent handoff acceptance	None	Agents act on handoff within their stated SLA

7. Work split: Jiv 15 points, Gabe 15 points

Sizing: S=1, M=2, L=3. Split follows the code boundary so the two streams rarely touch the same files.
Assumption: Gabe's strengths are not yet confirmed; swap items if needed but keep the point balance.

Jiv: agent, channels, compliance (backend core)

ID	Item	Pts	Files
J1	Real-estate agent v2: qualification state, prompt, stage detection (RE-01, RE-02)	3	agent/prompts/real_estate.j2 , agent/state.py , agent/stage_detection.py , agent/graph.py

ID	Item	Pts	Files
J2	Lead ingestion: Lead Ads leadgen fetch, CSV, dedupe (LEAD-01)	2	api/main.py webhooks, api/channels.py
J3	Tenant routing, Messenger and SMS hardening, window rules (CH-01, CH-02)	3	api/main.py , api/channels.py , api/db.py
J4	Handoff service: assignment, summary, push/SMS to agent (HAND-01 backend)	2	api/escalation_service.py , api/main.py
J5	Consent, notice, opt-out, audit log (CONS-01)	2	api/ , agent/prompts/base.j2
J6	Metrics events table and /metrics endpoint (MET-01 backend)	2	api/ , new migration
J7	Model bake-off on the Taglish eval set (Llama 3.3 70B vs alternatives)	1	evals/ , litellm_config.yaml

Gabe: surfaces, auth, data

ID	Item	Pts	Files
G1	Finish Neon+Clerk code: plan Tasks 7, 8, 10 (settings to /me/business , lib/api.ts bearer helper, delete supabase stub, onboarding signup, English copy)	3	mobile/ , frontend/src/OnboardingWizard.tsx
G2	Provision Neon and Clerk, deploy, cutover (plan Task 9, AUTH-01)	2	infra, env, scripts/
G3	Web inbox and conversation view with qualification card (UI-01)	3	frontend/src/dashboard/ , new inbox
G4	Agent roster import and assignment rule UI (HAND-02)	2	frontend/src/
G5	Pilot dashboard: reply-time AI vs human, funnel, per-agent table, CSV (MET-01 UI)	2	frontend/src/dashboard/
G6	Mobile: login, pipeline with new fields, push on handoff (UI-01)	3	mobile/app/

Contract between the two streams

Agree these in week 1 and freeze them in docs/api-contract.md : lead object (with the qualification fields), conversation message object, handoff payload, metrics event names, /leads , /metrics and /agents response shapes. Gabe builds against fixtures until the real endpoints land.

Dependencies

- G2 (deploy) before J3 can be tested against a real tenant.

- J6 feeds G5. J4 feeds the push in G6.
- J1 and J7 are independent of everything else; start there.

Proposed sequence (proposal, not researched)

Week	Jiv	Gabe
0	Discovery call with the Filipinohomes contact (section 9)	Freeze contract, start G1
1 to 2	J1, J7, J2	G1, G2
3 to 4	J3, J5, J6	G3, G4
5	J4	G5, G6
6	Pilot on a small agent group, shadow mode first (AI drafts, human sends)	Same

Shadow mode first is a deliberate safety choice: it produces reply-time and quality data before any real lead sees an AI message.

8. Decisions needed (with a default so work does not stall)

#	Decision	Default until decided	Why it is open
D1	Model for Taglish	Keep <code>groq/llama-3.3-70b-versatile</code> , bake off in J7	FilBench and Batayan are NLP benchmarks on older models, not sales dialogue (C6). SEA-LION scored about 11 points below GPT-4o. Test, do not assume.
D2	Text first, voice later	Text only	Voice was not researched.
D3	AI disclosure policy	Disclose on inbound (current), send privacy notice and opt-out in the first outbound message, never claim to be human	NPC Advisory 2024-04 requires transparency about AI processing, not a "bot must announce itself" line (C7). Meta and Viber bot rules are unresearched. This conflicts with the earlier "never reveal AI" preference, so Jiv decides after the compliance check.
D4	Reply channel priority	Messenger and SMS	Channel mix for PH buyers is [unverified] ; Filipinohomes contact will tell us.
D5	Tenant model	One business (Filipinhomes), many agents as handoff targets	Depends on whether each agent wants their own persona or number.
D6	Sending entity	Filipinhomes brand, sent through our accounts	Viber domestic rate needs a PH-headquartered sender (C5).

9. What still needs research

The first run answered 12 of roughly 80 claims. This is the remaining list, ordered by what blocks the pilot.

A. Ask the Filipinohomes contact (web search cannot answer these)

#	Question	Blocks
Q1	How do leads reach agents today (on-site message, email, Lead Ads, calls)? Can they expose a webhook, API or email forward?	J2, LEAD-01
Q2	Monthly lead volume, and real inquiry-to-first-reply times across agents. Is the "Inquiry Reply" metric computed by the platform?	Section 6 baseline
Q3	Which channels do their buyers actually use (Messenger, Viber, SMS, calls)?	D4
Q4	Who owns the Facebook Page and ad accounts, and can we get an app installed with <code>ADVERTISE</code> access (needs <code>lead_retrieval</code> , <code>pages_manage_metadata</code> and others, C8)?	J2
Q5	Which agents join the pilot, and how should assignment work?	HAND-02
Q6	Existing privacy notice and consent wording on their forms	CONS-01
Q7	Are leads handled on behalf of licensed brokers, and who supervises?	Section 9C
Q8	Commercial terms: pilot free, paid, revenue share, partner path	Business model

B. Run another deep-research pass (web can answer)

#	Topic	Blocks
R1	Meta Messenger policy for AI bots, 24-hour window and human-agent tag, outbound template rules	CH-02, D3
R2	Philippine SMS rules (NTC, telco registration, sender ID, Semaphore terms), consent for marketing SMS	CH-02, CONS-01
R3	Competitors and prior art: Lamudi, Dot Property and developer chat tools. No competitor price survived verification, and the Lamudi and botsatwork.ph price claims were refuted	Pricing, positioning
R4	PH buyer qualification facts (Pag-IBIG rules, pre-selling terms, OFW process) from primary sources. The blog figures in the research file are unverified	RE-01 content
R5	Current Taglish quality of Claude, Gemini, GPT and Llama on a sales-dialogue test, ideally our own eval	D1

C. Needs a lawyer or regulator source

- Does an AI qualifying and booking leads for licensed agents create PRC or DHSUD broker-licensing exposure? No usable source found (the DHSUD FAQ page was judged unreliable). Frame OttoBot as a tool acting for the licensed broker until told otherwise.

10. Risks

Risk	Impact	Mitigation
AI says something wrong about price or availability	Trust and legal	Agent never quotes a final price or availability; refers to the assigned agent. Shadow mode first.
Taglish quality below bar	Bad first impression	J7 bake-off, human handoff on low confidence, guardrails in <code>api/guardrails.py</code>
Messaging window or platform policy violation	Page or number banned	R1 and R2 before CH-02 goes live; block out-of-window sends
Single-tenant env config reaches production	Cross-tenant data leak	CH-01 is P0; add the two-tenant test
Lead Ads webhook design wrong	No leads arrive	J2 first, verified against a test form
Groq model IDs rot	Outage	Keep <code>_build_model_list</code> check; verify <code>/v1/models</code> before release (see FAILURES.md)
Partner expectations outrun a text-only MVP	Relationship	Section 1 non-goals shared with the contact in week 0

11. Docs to update when this ships

`CLAUDE.md` (current state), `README.md` (env and run steps), `.env.example`, `docs/api-contract.md` (new), `docs/design/guardrails.md` (any new screen patterns). Obsidian `Projects/ottobot/` gets a pointer to this file and the decisions D1 to D6, not a copy.